

Neural Networks

Data Mining
Prof. Dawn Woodard
School of ORIE
Cornell University

Outline

1 Training a Neural Network

2 Digit Recognition Example

Training a Neural Network

Basic idea of neural net training (via “Back-Propagation”):

- Start with some arbitrary values for the parameters α, β
- Calculate predictions \hat{y}_i for all observations i in the training data (“forward sweep”). Use these to calculate the error.
- Update the parameters by changing them in a way that decreases that error (“backward sweep”)
- Repeat this iterative process until:

Training a Neural Network

Continuous-outcome case:

Training error is measured using $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$ where n is the # training observations

Because each node is a differentiable function of its inputs, it is possible to calculate $\frac{\partial RSS}{\partial \alpha_{mj}}$ and $\frac{\partial RSS}{\partial \beta_{km}}$, the derivative of RSS w.r.t. the parameters

Training a Neural Network

Continuous-outcome case:

1. Start with arbitrary initial values

$\{\alpha_{mj}^{(0)}, \beta_{km}^{(0)} : m \in \{0, \dots, M\}, j \in \{0, \dots, p\}, k \in \{1, \dots, K\}\}$ of the parameters

2. Do the following update until the validation-set RMSE starts to increase:

$$\beta_{km}^{(r+1)} =$$

$$\alpha_{mj}^{(r+1)} =$$

Training a Neural Network

The amount by which the parameters are allowed to change in each iteration is controlled by the “learning rate” $\gamma \in (0, 1)$

Why don't we just keep iterating until the parameters stop changing (i.e. find a local minimum of RSS)?

Digit Recognition

Handwritten digits from U.S. postal envelopes (Example from Hastie, Tibshirani, and Friedman 2009):




FIGURE 1.2. *Examples of handwritten digits from U.S. postal envelopes.*

Digit Recognition

- Have a dataset with labeled digits. Goal is to learn a decision rule for identifying a digit from its image
- The predictors are the grayscale intensities of each of the pixels in the 16x16 image (256 continuous predictors that take values 0 to 255)
- Outcome variable is:



Digit Recognition

- This is a classification problem with $K = 10$ categories
- Such a decision rule would be used for automatic sorting of mail according to its destination
- Very high accuracy is required
- Some preprocessing has been done to align the digits spatially
- For illustration there are 320 digits in training set and 160 in test set

Digit Recognition

- To apply neural networks here we have to have 10 output nodes, one for each possible value of the outcome variable
- Various neural net structures were considered, including using two hidden layers that are “locally connected.” Locally connected means that not all possible edges are included in the graph; this makes the model less prone to overfitting
- Also the technique of “weight sharing” was used, meaning that some nodes in the hidden layer are forced to have the same parameters α (although they have slightly different inputs due to the local connectivity). Again this is done to avoid overfitting.

Digit Recognition

Five neural nets were trained on the training data:

- 1 No hidden layer, equivalent to “multinomial logistic regression” (logistic regression with >2 outcome values)
- 2 One hidden layer with 12 nodes
- 3 Two hidden layers locally connected
- 4 Two hidden layers, locally connected with weight sharing
- 5 Two hidden layers, locally connected, two levels of weight sharing

Digit Recognition

Results:




FIGURE 11.11. Test performance curves, as a function of the number of training epochs, for the five networks of Table 11.1 applied to the ZIP code data. (Le Cun, 1989)

Digit Recognition

Accuracy of 5 models on test data:

TABLE 11.1. *Test set performance of five different neural networks on a hand-written digit classification example (Le Cun, 1989).*

	Network Architecture	Links	Weights	% Correct
Net-1:	Single layer network	2570	2570	80.0%
Net-2:	Two layer network	3214	3214	87.0%
Net-3:	Locally connected	1226	1226	88.5%
Net-4:	Constrained network 1	2266	1132	94.0%
Net-5:	Constrained network 2	5194	1060	98.4%

Other Applications

The company “Alyuda” has created a commercial software product exclusively for fitting and applying neural networks.

Find their demo by googling “Alyuda”, and clicking on “NeuroIntelligence” and then “Take a Tour”